

Predicting Customer Churn in a Telecommunications Company

Problem Statement

Customer churn is a significant issue in the telecommunications industry. Retaining existing customers is often more cost-effective than acquiring new ones. By analyzing customer data, we can identify patterns and factors that contribute to churn, enabling the company to take proactive measures to retain customers.

1. Import Libraries

```
In [56]: # Data manipulation and analysis
import pandas as pd
import numpy as np

# Visualization
import matplotlib.pyplot as plt
import seaborn as sns

# Machine Learning
from sklearn.model_selection import train_test_split, GridSearchCV, StratifiedKFold
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.feature_selection import RFECV
from sklearn.metrics import classification_report, confusion_matrix, roc_curve
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier

# Handling imbalanced data
from imblearn.over_sampling import SMOTE

# Suppress warnings
import warnings

warnings.filterwarnings("ignore")
```

2. Load Dataset

```
In [57]: # Load the dataset
df = pd.read_csv("WA_Fn-UseC_-Telco-Customer-Churn.csv")

# Display first few rows
df.head()
```

Out [57]:

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLine
0	7590-VHVEG	Female	0	Yes	No	1	No	No phon servic
1	5575-GNVDE	Male	0	No	No	34	Yes	N
2	3668-QPYBK	Male	0	No	No	2	Yes	N
3	7795-CFOCW	Male	0	No	No	45	No	No phon servic
4	9237-HQITU	Female	0	No	No	2	Yes	N

5 rows x 21 columns

This dataset contains information about a telecom company's customers, including:

- Customer account information (customerID, tenure, contract type, payment method, monthly charges, total charges).
- Demographic information (gender, age range, senior citizen status, partner, dependents).
- Services subscribed (phone service, multiple lines, internet service, online security, online backup, device protection, tech support, streaming TV, streaming movies).
- The target variable is Churn, indicating whether the customer left within the last month.

In [58]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customerID            7043 non-null   object
1   gender                7043 non-null   object
2   SeniorCitizen         7043 non-null   int64
3   Partner               7043 non-null   object
4   Dependents            7043 non-null   object
5   tenure                7043 non-null   int64
6   PhoneService          7043 non-null   object
7   MultipleLines         7043 non-null   object
8   InternetService       7043 non-null   object
9   OnlineSecurity        7043 non-null   object
10  OnlineBackup          7043 non-null   object
11  DeviceProtection      7043 non-null   object
12  TechSupport           7043 non-null   object
13  StreamingTV           7043 non-null   object
14  StreamingMovies       7043 non-null   object
15  Contract              7043 non-null   object
16  PaperlessBilling      7043 non-null   object
17  PaymentMethod         7043 non-null   object
18  MonthlyCharges        7043 non-null   float64
19  TotalCharges          7043 non-null   object
20  Churn                 7043 non-null   object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

3. Data Cleaning for EDA

3.1. Convert TotalCharges to Numeric

```
In [59]: # Convert 'TotalCharges' to numeric, coerce errors to NaN
df["TotalCharges"] = pd.to_numeric(df["TotalCharges"], errors="coerce")

# Check for missing values in 'TotalCharges'
missing_total_charges = df["TotalCharges"].isnull().sum()
print(f"Total missing values in 'TotalCharges': {missing_total_charges}")
```

Total missing values in 'TotalCharges': 11

3.2 Handle Missing Values

```
In [60]: # Drop rows with missing 'TotalCharges'
df = df.dropna(subset=["TotalCharges"])

# Reset index
df.reset_index(drop=True, inplace=True)
```


3.3 Remove Unnecessary Columns

```
In [61]: # Drop 'customerID' as it's not useful for prediction
df = df.drop(columns=["customerID"])
```

4. Exploratory Data Analysis (EDA)

4.1 Churn Distribution

```
In [62]: # Count plot for 'Churn'
plt.figure(figsize=(6, 4))
sns.countplot(x="Churn", data=df, palette="Set2")
plt.title("Churn Distribution")
plt.xlabel("Churn")
plt.ylabel("Count")
plt.show()
```



The dataset is imbalanced with more 'No' churn instances than 'Yes'.

4.2 Churn Rate by Contract Type

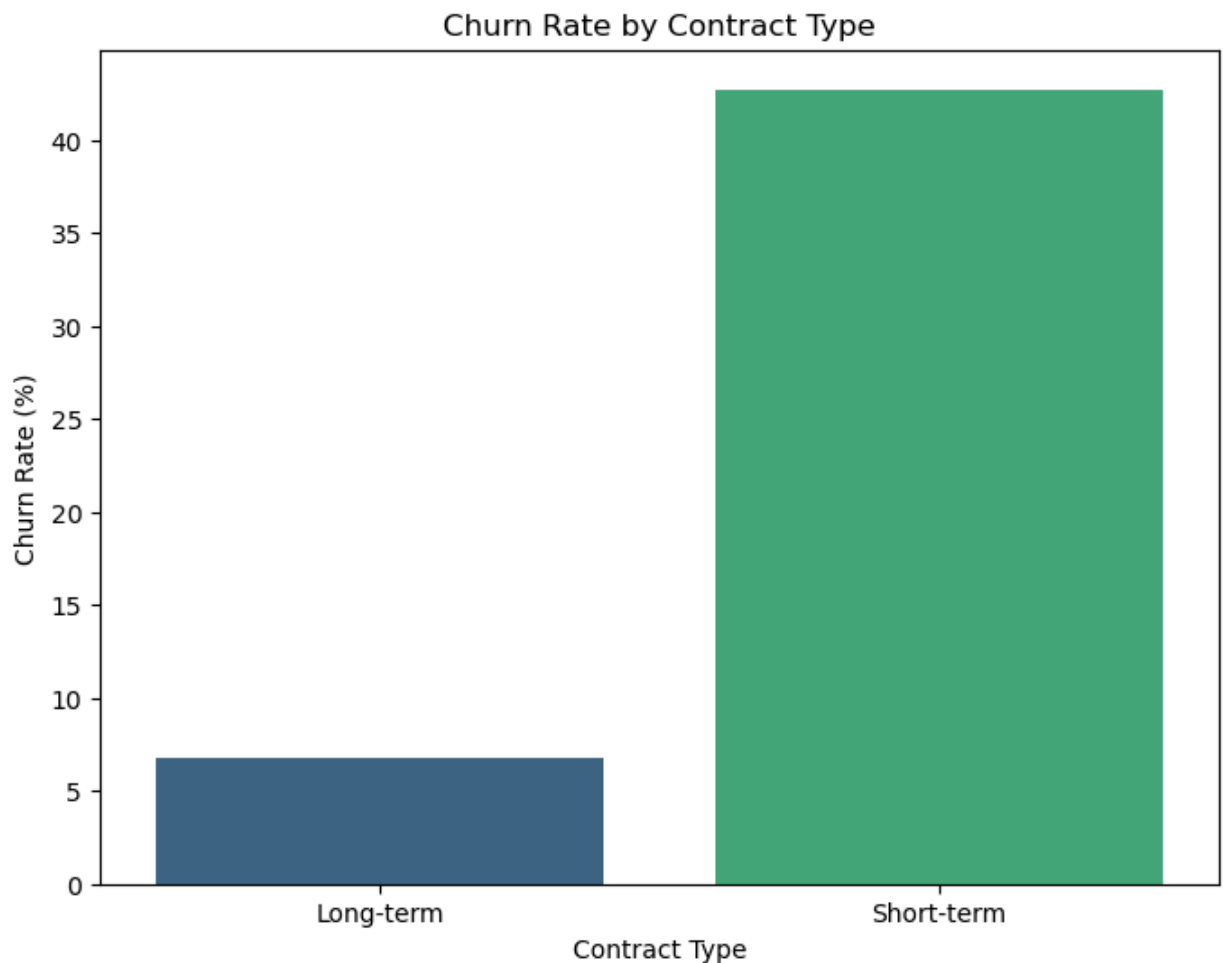
```
In [63]: # Create 'Contract_Type' variable
df["Contract_Type"] = df["Contract"].apply(
    lambda x: "Short-term" if x == "Month-to-month" else "Long-term"
)
# Map 'Churn' to numeric for calculation
df["Churn_numeric"] = df["Churn"].map({"No": 0, "Yes": 1})

# Calculate churn rate
churn_rate = df.groupby("Contract_Type")["Churn_numeric"].mean().reset_index()
churn_rate["Churn_Rate (%)"] = churn_rate["Churn_numeric"] * 100
churn_rate = churn_rate.drop(columns=["Churn_numeric"])
churn_rate
```

Out[63]:

	Contract_Type	Churn_Rate (%)
0	Long-term	6.778587
1	Short-term	42.709677

```
In [64]: plt.figure(figsize=(8, 6))
sns.barplot(x="Contract_Type", y="Churn_Rate (%)", data=churn_rate, palette="v")
plt.title("Churn Rate by Contract Type")
plt.xlabel("Contract Type")
plt.ylabel("Churn Rate (%)")
plt.show()
```

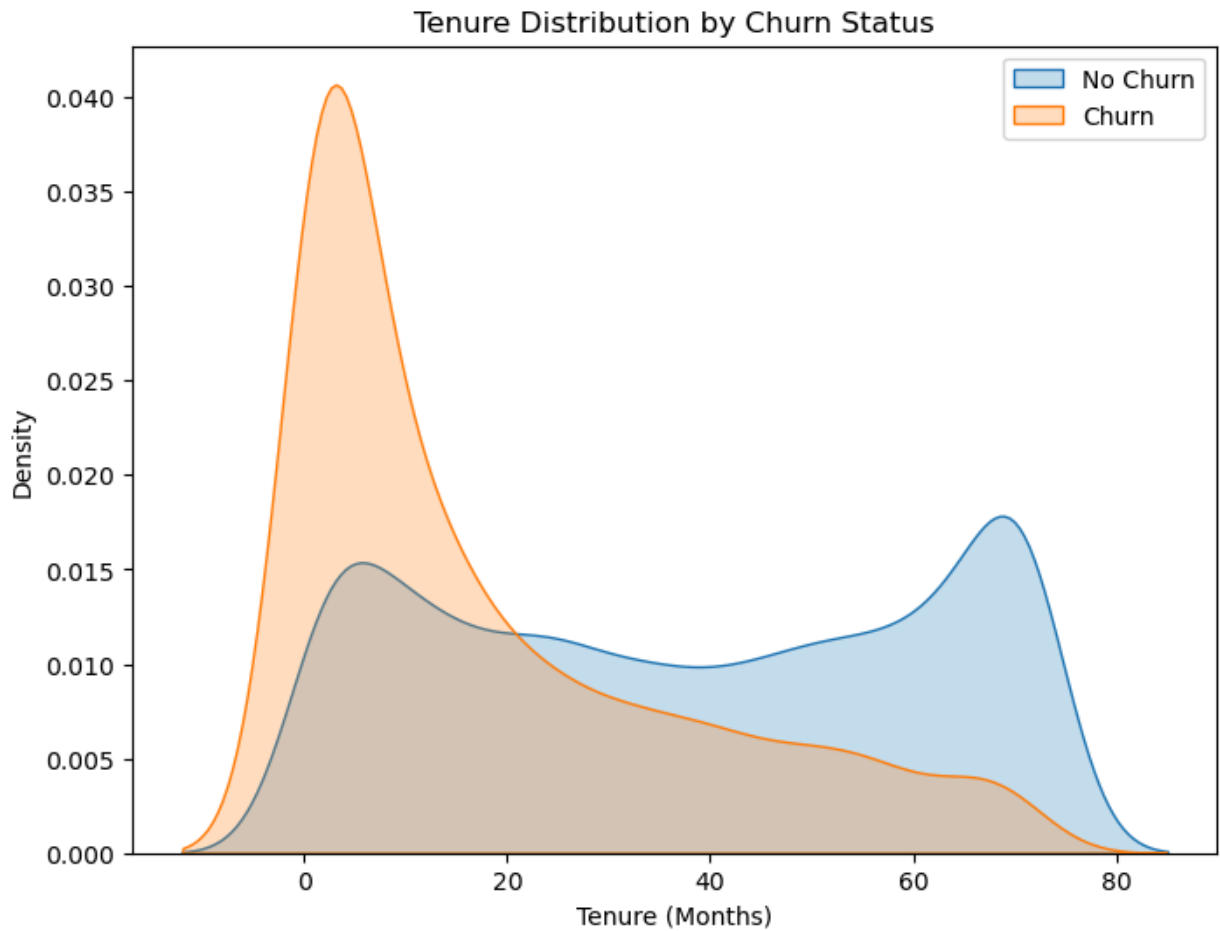


Short-term contracts have a significantly higher churn rate compared to long-term contracts.

4.3 Tenure vs. Churn

```
In [65]: # Distribution plots for tenure
plt.figure(figsize=(8, 6))
sns.kdeplot(df[df["Churn"] == "No"]["tenure"], label="No Churn", shade=True)
sns.kdeplot(df[df["Churn"] == "Yes"]["tenure"], label="Churn", shade=True)
plt.title("Tenure Distribution by Churn Status")
plt.xlabel("Tenure (Months)")
plt.ylabel("Density")
```


```
plt.legend()
plt.show()
```



Customers with shorter tenure are more likely to churn.

4.4 Monthly Charges vs. Churn

```
In [67]: # Distribution plots for MonthlyCharges
plt.figure(figsize=(8, 6))
sns.kdeplot(df[df["Churn"] == "No"]["MonthlyCharges"], label="No Churn", shade=True)
sns.kdeplot(df[df["Churn"] == "Yes"]["MonthlyCharges"], label="Churn", shade=True)
plt.title("Monthly Charges Distribution by Churn Status")
plt.xlabel("Monthly Charges ($)")
plt.ylabel("Density")
plt.legend()
plt.show()
```



Customers with higher monthly charges tend to churn more.

5. Data Preprocessing and Encoding

5.1 Encode Binary Variables

```
In [31]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7032 entries, 0 to 7031
Data columns (total 22 columns):
#   Column                Non-Null Count  Dtype
---  -
0   gender                 7032 non-null  object
1   SeniorCitizen          7032 non-null  int64
2   Partner                7032 non-null  object
3   Dependents             7032 non-null  object
4   tenure                 7032 non-null  int64
5   PhoneService           7032 non-null  object
6   MultipleLines           7032 non-null  object
7   InternetService        7032 non-null  object
8   OnlineSecurity         7032 non-null  object
9   OnlineBackup           7032 non-null  object
10  DeviceProtection       7032 non-null  object
11  TechSupport            7032 non-null  object
12  StreamingTV            7032 non-null  object
13  StreamingMovies        7032 non-null  object
14  Contract               7032 non-null  object
15  PaperlessBilling       7032 non-null  object
16  PaymentMethod          7032 non-null  object
17  MonthlyCharges         7032 non-null  float64
18  TotalCharges           7032 non-null  float64
19  Churn                  7032 non-null  object
20  Contract_Type          7032 non-null  object
21  Churn_numeric          7032 non-null  int64
dtypes: float64(2), int64(3), object(17)
memory usage: 1.2+ MB
```

```
In [32]: # List of binary variables to map
binary_features = ["Partner", "Dependents", "PhoneService", "PaperlessBilling"]

# Map 'Yes' to 1 and 'No' to 0
for col in binary_features:
    df[col] = df[col].map({"Yes": 1, "No": 0})
```

```
In [33]: df.head()
```

```
Out[33]:
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetSe
0	Female	0	1	0	1	0	No phone service	
1	Male	0	0	0	34	1	No	
2	Male	0	0	0	2	1	No	
3	Male	0	0	0	45	0	No phone service	
4	Female	0	0	0	2	1	No	Fiber

5 rows x 22 columns

```
In [34]: # Map 'Short-term' to 0 and 'Long-term' to 1
df["Contract_Type"] = df["Contract_Type"].map({"Short-term": 0, "Long-term": 1})
```


5.2 One-Hot Encode Categorical Variables

```
In [35]: # Identify categorical variables for one-hot encoding
categorical_vars = [
    "gender",
    "MultipleLines",
    "InternetService",
    "OnlineSecurity",
    "OnlineBackup",
    "DeviceProtection",
    "TechSupport",
    "StreamingTV",
    "StreamingMovies",
    "Contract",
    "PaymentMethod",
]

# Perform one-hot encoding
df_encoded = pd.get_dummies(df, columns=categorical_vars, drop_first=True)
```

```
In [36]: # Drop 'Contract' variable since we have created Contract_Type, we can drop the
df_encoded = df_encoded.drop(columns=["Contract_Two year", "Contract_One year"])
```

```
In [37]: df_encoded = df_encoded.drop(columns=["Churn_numeric"], axis=1)
```

5.3 Handle Remaining Missing Values

```
In [38]: df_encoded.head()
```

```
Out[38]:
```

	SeniorCitizen	Partner	Dependents	tenure	PhoneService	PaperlessBilling	MonthlyCharges
0	0	1	0	1	0	1	29.85
1	0	0	0	34	1	0	56.95
2	0	0	0	2	1	1	53.85
3	0	0	0	45	0	0	42.30
4	0	0	0	2	1	1	70.70

5 rows × 30 columns

```
In [39]: # Check for remaining missing values
df_encoded.isnull().sum()
```

```
Out[39]: SeniorCitizen      0
Partner      0
Dependents    0
tenure        0
PhoneService  0
PaperlessBilling  0
MonthlyCharges  0
TotalCharges  0
Churn         0
Contract_Type  0
gender_Male    0
MultipleLines_No phone service  0
MultipleLines_Yes  0
InternetService_Fiber optic  0
InternetService_No  0
OnlineSecurity_No internet service  0
OnlineSecurity_Yes  0
OnlineBackup_No internet service  0
OnlineBackup_Yes  0
DeviceProtection_No internet service  0
DeviceProtection_Yes  0
TechSupport_No internet service  0
TechSupport_Yes  0
StreamingTV_No internet service  0
StreamingTV_Yes  0
StreamingMovies_No internet service  0
StreamingMovies_Yes  0
PaymentMethod_Credit card (automatic)  0
PaymentMethod_Electronic check  0
PaymentMethod_Mailed check  0
dtype: int64
```

5.4 Feature Scaling

```
In [40]: # List of numerical features
numerical_features = ["tenure", "MonthlyCharges", "TotalCharges"]

# Initialize the scaler
scaler = StandardScaler()

# Scale numerical features
df_encoded[numerical_features] = scaler.fit_transform(df_encoded[numerical_fea
```

6. Machine Learning Modeling

6.1 Define Features and Target

```
In [41]: # Define target variable 'Churn'
y = df_encoded["Churn"]

# Define feature set 'X'
X = df_encoded.drop(columns=["Churn"])
```

6.2 Handling Class Imbalance

```
In [42]: # Check class distribution
y.value_counts(normalize=True)
```

```
Out[42]: Churn
0      0.734215
1      0.265785
Name: proportion, dtype: float64
```

the dataset is imbalanced. So applying SMOTE to balance classes

```
In [43]: # Initialize SMOTE
smote = SMOTE(random_state=42)

# Fit SMOTE to features and target
X_resampled, y_resampled = smote.fit_resample(X, y)

# Check new class distribution
y_resampled.value_counts(normalize=True)
```

```
Out[43]: Churn
0      0.5
1      0.5
Name: proportion, dtype: float64
```

6.3 Split the Data

```
In [44]: # Split into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(
    X_resampled, y_resampled, test_size=0.2, random_state=42, stratify=y_resampled
)
```

6.4 Feature Selection

Use Recursive Feature Elimination with Cross-Validation (RFECV)

```
In [45]: # Initialize Logistic Regression model
model = LogisticRegression(max_iter=1000, random_state=42)

# Initialize RFECV
rfecv = RFECV(estimator=model, step=1, cv=StratifiedKFold(5), scoring="accuracy")

# Fit RFECV
rfecv.fit(X_train, y_train)

# Optimal number of features
print(f"Optimal number of features: {rfecv.n_features_}")

# Selected features
selected_features = X_train.columns[rfecv.support_]
print("Selected features:", list(selected_features))
```

Optimal number of features: 25

Selected features: ['Dependents', 'tenure', 'PhoneService', 'MonthlyCharges', 'TotalCharges', 'Contract_Type', 'MultipleLines_No phone service', 'MultipleLines_Yes', 'InternetService_Fiber optic', 'InternetService_No', 'OnlineSecurity_No internet service', 'OnlineSecurity_Yes', 'OnlineBackup_No internet service', 'OnlineBackup_Yes', 'DeviceProtection_No internet service', 'DeviceProtection_Yes', 'TechSupport_No internet service', 'TechSupport_Yes', 'StreamingTV_No internet service', 'StreamingTV_Yes', 'StreamingMovies_No internet service', 'StreamingMovies_Yes', 'PaymentMethod_Credit card (automatic)', 'PaymentMethod_Electronic check', 'PaymentMethod_Mailed check']

6.5 Model Training and Evaluation

6.5.1 Logistic Regression

```
In [46]: # Create Logistic Regression model with selected features
lr_model = LogisticRegression(max_iter=1000, random_state=42)
lr_model.fit(X_train[selected_features], y_train)


# Predict on test set
y_pred_lr = lr_model.predict(X_test[selected_features])

# Classification report
print("Logistic Regression Classification Report:")
print(classification_report(y_test, y_pred_lr))
```

Logistic Regression Classification Report:


	precision	recall	f1-score	support
0	0.82	0.74	0.78	1033
1	0.76	0.84	0.80	1033
accuracy			0.79	2066
macro avg	0.79	0.79	0.79	2066
weighted avg	0.79	0.79	0.79	2066

```
In [47]: # Confusion matrix
cm = confusion_matrix(y_test, y_pred_lr)
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues")
plt.title("Logistic Regression Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```



```
In [48]: # ROC Curve
y_pred_proba = lr_model.predict_proba(X_test[selected_features])[:, 1]
fpr, tpr, _ = roc_curve(y_test, y_pred_proba)
auc_score = auc(fpr, tpr)

plt.figure()
plt.plot(fpr, tpr, label=f"AUC = {auc_score:.2f}")
plt.plot([0, 1], [0, 1], linestyle="--")
plt.title("ROC Curve - Logistic Regression")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend()
plt.show()
```



6.5.2 Random Forest Classifier

```
In [49]: # Define parameter grid
param_grid = {
    "n_estimators": [100, 200],
    "max_depth": [None, 10, 20],
    "min_samples_split": [2, 5],
    "min_samples_leaf": [1, 2],
}

# Initialize Random Forest model
rf = RandomForestClassifier(random_state=42)

# Initialize GridSearchCV
grid_search = GridSearchCV(
    estimator=rf, param_grid=param_grid, cv=5, scoring="accuracy", n_jobs=-1
)

# Fit GridSearchCV
grid_search.fit(X_train[selected_features], y_train)

# Best parameters
print(f"Best Parameters: {grid_search.best_params}")
```

Best Parameters: {'max_depth': None, 'min_samples_leaf': 2, 'min_samples_split': 2, 'n_estimators': 200}

```
In [50]: # Best Random Forest model
best_rf = grid_search.best_estimator_
best_rf.fit(X_train[selected_features], y_train)
```


```
# Predict on test set
y_pred_rf = best_rf.predict(X_test[selected_features])

# Classification report
print("Random Forest Classification Report:")
print(classification_report(y_test, y_pred_rf))
```

Random Forest Classification Report:


	precision	recall	f1-score	support
0	0.85	0.78	0.81	1033
1	0.80	0.86	0.83	1033
accuracy			0.82	2066
macro avg	0.82	0.82	0.82	2066
weighted avg	0.82	0.82	0.82	2066

```
In [51]: # Confusion matrix
cm_rf = confusion_matrix(y_test, y_pred_rf)
sns.heatmap(cm_rf, annot=True, fmt="d", cmap="Greens")
plt.title("Random Forest Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```



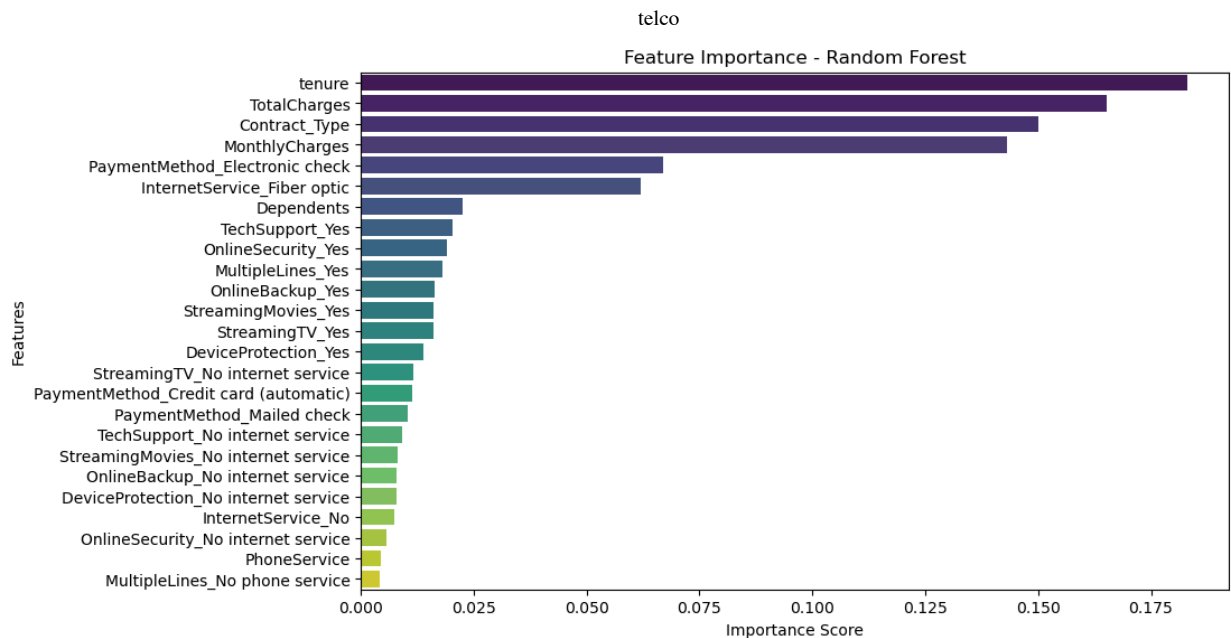
```
In [52]: # ROC Curve
y_pred_proba_rf = best_rf.predict_proba(X_test[selected_features])[:, 1]
fpr_rf, tpr_rf, _ = roc_curve(y_test, y_pred_proba_rf)
auc_score_rf = auc(fpr_rf, tpr_rf)
```

```
plt.figure()
plt.plot(fpr_rf, tpr_rf, label=f"AUC = {auc_score_rf:.2f}", color="green")
plt.plot([0, 1], [0, 1], linestyle="--")
plt.title("ROC Curve - Random Forest")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.legend()
plt.show()
```



```
In [53]: # Feature importance
importances = best_rf.feature_importances_
feature_importance = pd.Series(importances, index=selected_features).sort_values(
    ascending=False
)

plt.figure(figsize=(10, 6))
sns.barplot(x=feature_importance, y=feature_importance.index, palette="viridis")
plt.title("Feature Importance - Random Forest")
plt.xlabel("Importance Score")
plt.ylabel("Features")
plt.show()
```

7. Conclusion

Observation:

- Model performance:
 - Logistic Regression:
 - Accuracy: 0.79
 - AUC: 0.88
 - Random Forest:
 - Accuracy: 0.82
 - AUC: 0.89
 - Both Logistic regression and Random Forest models perform well after handling class imbalance and feature selection.
 - Random Forest typically has higher accuracy and better ability to capture complex patterns.
- key predictors of Churn:
 - Tenure: Lower tenure increases the likelihood of churn.
 - Contract_Type: Short-term contracts are associated with higher churn.
 - MonthlyCharges: Higher charges are linked to increased churn.

Recommendations:

- Customer Retention Strategies:
 - Tenure
 - Implement loyalty programs targeting new customers to increase tenure.
 - Contract
 - Encourage customers to switch to long-term contracts through incentives.
 - Provide discounts for customers who switch to longer-term contracts.

- Create attractive packages that encourage customers to opt for longer contracts.
- Monthly Chargers
 - Re-evaluate pricing models to ensure competitiveness.